

***Show the inferred beta vectors and indicate how they map to the true topics above.***

The true topics distribution (beta) is below figure:

```
vocabulary = [  
    "bass",  
    "pike",  
    "deep",  
    "tuba",  
    "horn",  
    "catapult",  
]  
  
beta = np.array(  
    [  
        [0.4, 0.4, 0.2, 0.0, 0.0, 0.0],  
        [0.0, 0.3, 0.1, 0.0, 0.3, 0.3],  
        [0.3, 0.0, 0.2, 0.3, 0.2, 0.0],  
    ]  
)
```

The inferred beta output is below:

```
Topic 0:  
bass: 0.2972  
pike: 0.2451  
deep: 0.1667  
tuba: 0.1074  
horn: 0.1017  
catapult: 0.0820  
Topic 1:  
bass: 0.2274  
pike: 0.1737  
deep: 0.2284  
tuba: 0.1655  
horn: 0.1678  
catapult: 0.0371  
Topic 2:  
bass: 0.2091  
pike: 0.2780  
deep: 0.1240  
tuba: 0.0800  
horn: 0.1835  
catapult: 0.1254
```

As the training process is unsupervised learning, the topic orders of inferred beta are not necessarily the same with the topic orders of input beta. From the above specific output, we could tell that the inferred beta's topic 0 matched relatively well with the input beta's topic 0 (first row) – with “bass” and “pike” being the highest frequency, and “tuba”, “horn”, and “catapult” being the lowest frequency.

For the inferred beta's topic 1, it seemed that the best match is the input beta's topic 2 (third row). However, “pike” should have lower frequency and “tuba” should have higher frequency, and thus not a complete match.

For the inferred beta's topic 2, it seemed that the best match is the input beta's topic 1 (second row). However, “bass” should have lower frequency and “horn” and “catapult” should have higher frequency, and thus not a complete match.

I think the key reasons for the mismatch between input beta and inferred beta are because the corpus is relatively small and training data is limited, leading to limited accuracy of output. If we could enlarge the training data, I believe the trained output would be closer to the true topics distribution.